Dhwani Jain

Ghaziabad, India | +91-8868032126 | dhwanijain731@gmail.com | linkedin.com/in/dhwani-jain | github.com/heydhwani

About

AI/ML Engineer with hands-on experience in building end-to-end Machine Learning pipelines and Generative AI applications using LLMs, RAG, and LangChain. Currently exploring Agentic AI — designing autonomous, multi-step AI systems that go beyond chat. I don't just learn AI, I build with it.

Education

Ajay Kumar Garg Engineering College — Ghaziabad, India
Bachelor of Technology in Computer Science | CGPA: 8.1 — *Oct 2023 – Sept 2027*

Experience

Computer Society of India — Ghaziabad, Uttar Pradesh
Machine Learning Developer — *Oct 2024 – Dec 2024*

- Enhanced machine learning model performance by 10% through feature engineering, hyperparameter tuning, and systematic model evaluation.
- Collaborated with peers to review model outputs, debug ML pipelines, and improve overall system reliability while meeting 100% of assigned deadlines.

Beyond Chatbots (Instagram) — Online
AI Content Creator — *May 2026 - Present*

- Built a public facing AI solutions platform where real world user problems are solved live using Generative AI tools, LLMs, and automation workflows.
- Translated complex AI concepts into accessible, solution-driven content for a growing tech audience.

Technical Skills

Programming Languages: Python, C, C++
Machine Learning: Supervised Learning, Classification, Regression, Random Forest, Gradient Boosting, XGBoost, Model Training, Model Evaluation, Model Fine-tuning
Deep Learning: Neural Networks, CNNs, RNNs, LSTMs, Transformers, Attention Mechanism
Generative AI: Large Language Models (LLMs), Prompt Engineering, RAG (Retrieval-Augmented Generation), LangChain, Vector Databases, Embeddings, Hugging Face Transformers
Data Analysis: Data Preprocessing, Exploratory Data Analysis (EDA), Feature Engineering, Data Cleaning
Libraries & Frameworks: Scikit-learn, Pandas, NumPy, Matplotlib, TensorFlow/PyTorch (basics)
Web & Deployment: REST APIs, FastAPI, Streamlit, Gradio, Model Deployment, API Integration
Tools & Platforms: Git, GitHub, Postman, VS Code, Jupyter Notebook, Google Colab, FAISS (Vector DB)

Projects

GenAI PDF Chatbot | *LangChain, FAISS, HuggingFace, Gemini 2.0, Streamlit* | GitHub | Live Demo

- Developed an end-to-end Retrieval-Augmented Generation (RAG) pipeline using LangChain and Google Gemini 2.0 Flash Lite LLM to enable natural language question-answering over uploaded PDF documents.
- Implemented semantic document chunking using RecursiveCharacterTextSplitter (chunk size: 500, overlap: 100) and generated vector embeddings via HuggingFace sentence-transformers (all-MiniLM-L6-v2), stored in a FAISS vector database for similarity-based retrieval.
- Engineered a custom prompt template ensuring hallucination-free, context-grounded responses by strictly limiting LLM output to retrieved document context.
- Deployed an interactive Streamlit web application with real-time PDF upload, session-based query handling, and example prompt suggestions, accessible via public URL.

AI Critical Minutes | *Random Forest, FastAPI, Streamlit* | GitHub | Live Demo

- Architected an end-to-end AI emergency decision support system using Random Forest multi-class classification to identify 3 emergency risk levels from 8 vital parameters, achieving 97% accuracy with high recall for critical cases.
- Generated a custom synthetic dataset of 5,000 samples using medically realistic ranges and weighted sampling to model multi-signal physiological instability.

- Implemented an AI decision engine that translates model predictions into structured 2+ Do and Don't emergency guidance, improving response clarity; integrated a Streamlit frontend with real-time voice alerts.

Customer Churn Prediction | *FastAPI, Gradient Boosting Classifier, Scikit-learn, Streamlit* | GitHub | Live Demo

- Engineered and deployed a customer churn prediction system using a Gradient Boosting Classifier on 11 input features, achieving approximately 87% accuracy.
- Constructed an end-to-end machine learning pipeline on 10,000+ customer records, enabling data-driven churn analysis and retention insights.
- Evaluated model performance using 3 key metrics (accuracy, precision, recall) to ensure balanced churn vs. non-churn predictions across unseen data.

Achievements

Participated in the **Flipkart Runway Season 5** Online Assessment, gaining exposure to industry-level problem solving scenarios.

Ranked among the **Top 15 participants** in a college-level Coding Quest competitive programming event.